

Filtered Intelligence Signals - Run 74edfc1f

233 High-Quality Signals

1. Databricks hits \$188B valuation, extending its run as AI's favorite second act

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Databricks hits \$188B valuation, extending its run as AI's favorite second act. Databricks has remade its image into an AI company and has published research on the cost savings of open weight AI models for coding.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/databricks-hits-188b-valuation-extending-its-run-as-ais-favorite-second-act>

2. Lark Health, Samsung team up on AI-powered health coach for U.S. seniors - Fierce Healthcare

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Lark Health, Samsung team up on AI-powered health coach for U.S. seniors - Fierce Healthcare. Lark Health, Samsung team up on AI-powered health coach for U.S. seniors Fierce Healthcare

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.fiercehealthcare.com/ai-and-machine-learning/lark-health-samsung-team-ai-powered-health-coach-us-seniors>

3. Answer-Conditioned Chains of Thought Degrade Verifiable-Reasoning Distillation in Large Language Models

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Answer-Conditioned Chains of Thought Degrade Verifiable-Reasoning Distillation in Large Language Models. A standard recipe for distilling the reasoning ability of large language models (LLMs) is to sample chains of thought from the model, keep those that reach the correct final answer, and fine-tune on the survivors. When sampling fails, a common fix shows the generator the gold answer and asks it to write a chain that reaches that answer. We show that this second step degrades the training data

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14552v1>

4. Latent Trajectory Discrimination for AI-Generated Text Detection

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Latent Trajectory Discrimination for AI-Generated Text Detection. Most existing approaches to AI-Generated Text Detection (AIGTD) treat documents as static objects and base their decisions on aggregate statistics or globally compressed embeddings. However, this perspective overlooks the inherently dynamic nature of autoregressive generation, where content evolves progressively through the latent space. In this paper, we reformulate AIGTD as the problem of distinguishing between latent generation

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14967v1>

5. GM's AI tools could cut the car development timeline in half

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: GM's AI tools could cut the car development timeline in half A great acceleration is coming to the auto industry thanks to artificial intelligence, and one of the giants of the carmaking world is using the technology to reinvent its product pipeline. In recent years, General Motors has dabbled with AI and virtual si... A great acceleration is coming to the auto industry thanks to artificial intelligence, and one of the giants of the carmaking world is using the technology to reinvent its product p

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://www.fastcompany.com/91571405/gm-ai-tools-could-cut-the-car-development-timeline-in-half>

6. CRISPR gets a power boost from AI-designed 'molecular scissors'

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: CRISPR gets a power boost from AI-designed 'molecular scissors' Researchers used artificial intelligence to design functional CRISPR enzymes not seen in nature. An enzyme cuts DNA (artists impression). Synthetic protein-cleaving enzymes were designed quickly with the help of artificial intelligence.Credit: Artur Plawgo/Science Photo Library Scientists have ... [+2866 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nature.com/articles/d41586-026-02217-w>

7. Prepare These 5 Assets Before Your AI Agents Take On More Work

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Prepare These 5 Assets Before Your AI Agents Take On More Work. How to define recurring work, give AI the right context, explain what high-quality work looks like, and decide where human judgment is still needed. The post Prepare These 5 Assets Before Your AI Agents Take On More Work appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://towardsdatascience.com/prepare-these-5-assets-before-your-ai-agents-take-on-more-work>

8. How to Get the Most Out of Claude Fable 5

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: How to Get the Most Out of Claude Fable 5. Maximize your Claude Fable 5 usage The post How to Get the Most Out of Claude Fable 5 appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/how-to-get-the-most-out-of-claude-fable-5>

9. The Misclassification of Autistic Writing as AI-Generated

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: The Misclassification of Autistic Writing as AI-Generated. Recent findings suggest that detection models for artificial intelligence (AI) cannot accurately identify AI-generated text and may exhibit bias against certain minority groups. In the present study, anecdotal claims that autistic writers more often have their work flagged as AI-generated are examined empirically. A corpus of approximately 60,000 Reddit posts split into "likely-autistic" and "general-Reddit" subcorpora is used to compare

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14729v1>

10. Context Engineering for RAG Question Parsing: From a Raw Question to Typed Fields That Steer Retrieval and Generation

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Context Engineering for RAG Question Parsing: From a Raw Question to Typed Fields That Steer Retrieval and Generation. Enterprise Document Intelligence [Vol.1 #6quarter] - Question parsing takes one messy string and writes four typed pieces, each read by a different downstream call The post Context Engineering for RAG Question Parsing: From a Raw Question to Typed Fields That Steer Retrieval and Generation appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/context-engineering-for-rag-question-parsing-from-a-raw-question-to-typed-fields-that-steer-retrieval-and-generation>

11. US Geological Survey seeks enhanced AI function to watch for earthquakes - Washington Technology

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: US Geological Survey seeks enhanced AI function to watch for earthquakes - Washington Technology. US Geological Survey seeks enhanced AI function to watch for earthquakes Washington Technology

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.washingtontechnology.com/contracts/2026/07/us-geological-survey-seeks-enhanced-ai-function-watch-earthquakes/414779>

12. Charting a Responsible Course for AI - Pharmacy Times

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Charting a Responsible Course for AI - Pharmacy Times. Charting a Responsible Course for AI Pharmacy Times

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.pharmacytimes.com/view/charting-a-responsible-course-for-ai>

13. Using AI for FOIA in 2026: What requesters should know - MuckRock

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Using AI for FOIA in 2026: What requesters should know - MuckRock. Using AI for FOIA in 2026: What requesters should know MuckRock

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.muckrock.com/news/archives/2026/jul/15/using-ai-for-foia-in-2026-what-requesters-should-know>

14. AI Didn't Make Programming Easier. It Just Made It Differently Difficult - Communications of the ACM

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI Didn't Make Programming Easier. It Just Made It Differently Difficult - Communications of the ACM. AI Didn't Make Programming Easier. It Just Made It Differently Difficult Communications of the ACM

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://cacm.acm.org/opinion/ai-didnt-make-programming-easier-it-just-made-it-differently-difficult>

15. How I Mastered Data Structures and Algorithms for ML (In 6 Weeks)

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: How I Mastered Data Structures and Algorithms for ML (In 6 Weeks). The strategies, questions, and process I used to ace coding interviews. The post How I Mastered Data Structures and Algorithms for ML (In 6 Weeks) appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/how-i-mastered-data-structures-algorithms-for-ml-in-6-weeks>

16. Most RAG Hallucinations Are Retrieval Failures: How the Retrieval Brick Decides What the Model Can Invent

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Most RAG Hallucinations Are Retrieval Failures: How the Retrieval Brick Decides What the Model Can Invent. Enterprise Document Intelligence [Vol.1 #7quinquies] - Hallucination is usually garbage-in. Fix retrieval, and the model has nothing left to make up The post Most RAG Hallucinations Are Retrieval Failures: How the Retrieval Brick Decides What the Model Can Invent appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/most-rag-hallucinations-are-retrieval-failures-how-the-retrieval-brick-decides-what-the-model-can-invent>

17. Alan Turing's biggest AI assumption may have been wrong - ScienceDaily

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Alan Turing's biggest AI assumption may have been wrong - ScienceDaily. Alan Turing's biggest AI assumption may have been wrong ScienceDaily

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.sciencedaily.com/releases/2026/07/260713084850.htm>

18. AI agents create virtual playgrounds to help robots get crucial training data - MIT News

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AI agents create virtual playgrounds to help robots get crucial training data - MIT News. AI agents create virtual playgrounds to help robots get crucial training data MIT News

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://news.mit.edu/2026/ai-agents-create-virtual-playgrounds-to-help-robots-get-crucial-training-data-0713>

19. Pydantic + OpenAI: The Cleanest Way to Get Structured Outputs from LLMs

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Pydantic + OpenAI: The Cleanest Way to Get Structured Outputs from LLMs. How to stop parsing JSON by hand and start trusting your model's output The post Pydantic + OpenAI: The Cleanest Way to Get Structured Outputs from LLMs appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://towardsdatascience.com/pydantic-openai-the-cleanest-way-to-get-structured-outputs-from-llms>

20. How I'm Making Sure My Analytics Career Doesn't Get Eaten by AI

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: How I'm Making Sure My Analytics Career Doesn't Get Eaten by AI. The analytics career I signed up for five years ago doesn't exist anymore, and honestly, I am fine with that. The post How I'm Making Sure My Analytics Career Doesn't Get Eaten by AI appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/how-im-making-sure-my-analytics-career-doesnt-get-eaten-by-ai>

21. ClawTeams

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: ClawTeams. The first goal-driven, proactive AI team for e-commerce. ClawTeams is an AI employee platform for e-commerce sellers. Instead of hiring specialists-or doing everything yourself-you get a coordinated AI team that thinks, plans, and executes like real employees. One goal. One team. Zero micromanagement. Tell your team lead what you want-"Increase Q4 revenue by 20%"-and they break it down, assign specialists, and run the plan. You get updates in Slack or Discord. High-stakes decisions wa

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/clawteams>

22. How Much Does It Actually Cost to Run a Local LLM? (Euros per Million Tokens, Measured)

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: How Much Does It Actually Cost to Run a Local LLM? (Euros per Million Tokens, Measured). I measured the actual GPU electricity for eight local models on one RTX 3090 - and the cheapest wasn't the smallest, nor the priciest the biggest. The post How Much Does It Actually Cost to Run a Local LLM? (Euros per Million Tokens, Measured) appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/how-much-does-it-actually-cost-to-run-a-local-llm-e-per-million-tokens-measured>

23. Agentic RAG: Let the Agent Search

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Agentic RAG: Let the Agent Search. A minimal OpenAI Agents SDK implementation where retrieval becomes a search-read-decide loop The post Agentic RAG: Let the Agent Search appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/agentic-rag-let-the-agent-search>

24. Building Models in Two Worlds: From Latent Constructs to Behavioral Signals

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Building Models in Two Worlds: From Latent Constructs to Behavioral Signals. My PhD models tried to explain why people engage. My industry models predict who will. The statistics barely changed. Everything around them did. The post Building Models in Two Worlds: From Latent Constructs to Behavioral Signals appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/building-models-in-two-worlds-from-latent-constructs-to-behavioral-signals>

25. Apple Watch Accounts for 90% of AI Smartwatch Shipments

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Apple Watch Accounts for 90% of AI Smartwatch Shipments Apple accounted for roughly 90% of all Edge AI-capable smartwatch shipments in the first quarter of 2026, according to new data from Counterpoint Research. That dominance came as Edge AI penetration across the broader smartwatch market grew 70% year-over... Apple accounted for roughly 90% of all Edge AI-capable smartwatch shipments in the first quarter of 2026, according to new data from Counterpoint Research. That dominance came as Edge AI p

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.macrumors.com/2026/07/10/apple-watch-accounts-for-most-ai-smartwatches>

26. Happy Birthday, AI-and Happy Birthday to You Too, Psychology

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Happy Birthday, AI-and Happy Birthday to You Too, Psychology While you're putting candles on your cake to celebrate the 70th birthday of artificial intelligence, don't forget to also mark the

70th birthday of cognitive science. June 18, 1956 was a memorable day. It was the day the term artificial intelligence was coined at the start of the Dartmouth Summer Research Project on Artificial Intelligence at Dartmouth College in ... [+5185 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.psychologytoday.com/us/blog/keeping-those-words-in-mind/202607/happy-birthday-ai-and-happy-birthday-to-you-too-psychology>

27. Daily Deal: The Premium Machine Learning Artificial Intelligence Super Bundle

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Daily Deal: The Premium Machine Learning Artificial Intelligence Super Bundle The Premium Machine Learning Artificial Intelligence Super Bundle is a twelve-course digital toolkit designed to take you from coding novice to AI architect. The curriculum breaks down complex data science concepts into actionable steps, guiding you through h... from the good-deals-on-cool-stuff dept The Premium Machine Learning Artificial Intelligence Super Bundle is a twelve-course digital toolkit designed to take you

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <http://www.techdirt.com/2026/07/08/daily-deal-the-premium-machine-learning-artificial-intelligenc-e-super-bundle>

28. ExploreYC

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: ExploreYC. Open-source API for Y Combinator & a16z company data. One open-source API for startup data across Y Combinator AND a16z - 6,600+ companies with funding, stage, IPO/M&A exits, and founders. Filter by VC (yc/a16z/all), batch, industry, country, or search. Grab a free API key, read the docs (curl/Node/Python + Swagger), and build in 30 seconds. Plus the full web app: map, analytics, funding data, a live hiring board, and AI tools.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/yc-company-explorer>

29. New AI model reveals how neutron star mergers forge heavy elements

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: New AI model reveals how neutron star mergers forge heavy elements Researchers have created an AI-based simulation that makes it much faster to model how neutron star mergers produce many of the universe's heaviest elements. The new tool could improve predictions of these powerful explosions while helping scientists bet... Researchers have developed a new artificial intelligence powered simulation that could significantly improve our understanding of how the universe creates many of its heavi

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.sciencedaily.com/releases/2026/06/260626030426.htm>

30. Uber salaries revealed: See how much the ride-hailing company pays product managers, software engineers, and others

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Uber salaries revealed: See how much the ride-hailing company pays product managers, software engineers, and others Uber pays many data scientists, software engineers, and other jobs close to \$300,000 a year, while some roles make even more, H-1B visa data shows. Uber has filed fewer H-1B visa applications so far this year, according to federal data. Bloomberg/Getty Images Uber is hiring in the US, though it's pulled back on visa applications for fore... [+3961 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.businessinsider.com/uber-salary-data-revealed-how-much-workers-make-2026-7>

31. AI is 'not smart' so what's next in artificial intelligence?

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AI is 'not smart' so what's next in artificial intelligence? Leading AI researcher Yan LeCun has a start-up which is developing a more flexible AI system. "We don't have robots that are nearly as good at understanding the physical world as a rat," says Yann LeCun, one of the leading figures in the world of artificial intelligence. He worked at Faceboo... [+6605 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.bbc.com/news/articles/cj6gr0xkyr3o?xtor=AL-72-%5Bpartner%5D-%5Byahoo.north.america%5D-%5Bheadline%5D-%5Bnews%5D-%5Bbizdev%5D-%5Bisapi%5D>

32. 10 YouTube Channels Keeping You Ahead in AI

Source: KDnuggets (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: 10 YouTube Channels Keeping You Ahead in AI. Explore 10 YouTube channels for AI engineers covering paper breakdowns, coding tutorials, and industry analysis.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.kdnuggets.com/10-youtube-channels-keeping-you-ahead-in-ai>

33. Tencent EdgeOne Makers

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Tencent EdgeOne Makers. Ship AI agents like web apps, in minutes.. Tencent EdgeOne Makers is an edge platform for modern web apps and AI agents. Build with your preferred frameworks and deploy through familiar CLI, Git, and CI/CD workflows. Get built-in agent runtime, sandboxed tools, memory, observability, model gateway support, serverless functions, and storage-without stitching together complex infrastructure. Add AI agents to existing products or launch new AI applications in minutes. Deploy

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/tencent-edgeone-2/launches/tencent-edgeone-makers>

34. AgentX

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AgentX. Evaluate AI agent, pinpoint issues, and fix with one click.. Evaluate AI agents before they fail. Create test suites, run evaluations, and pinpoint issues before they reach production. AgentX provides full observability and traceability for your AI agents. AI analysis not only identifies problems but also suggests fixes-like an AI doctor for your agents. Simulate run your agents across multiple LLM providers to compare performance, cost, and latency, helping you make better decisions abo

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/agentx/launches/agentx-4>

35. These AI Scams All Have Red Flags. Here's How to Spot Them

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: These AI Scams All Have Red Flags. Here's How to Spot Them The best thing you can do is to stay cautious and never hand over sensitive data. In the past couple of years, I've noticed an uptick in scam emails that get past Gmail's spam filters. Since I'm an author, the emails I get are usually from so-called book clubs, filled with empty p... [+7694 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.cnet.com/tech/services-and-software/ai-scams-have-red-flags-how-to-spot-them>

36. Import AI 460: Reward hacking society, RSI data from Anthropic; and RL-based quadcopter racing

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Import AI 460: Reward hacking society, RSI data from Anthropic; and RL-based quadcopter racing. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now Society can be reward-hacked, just like cyber environments:...Imagine an army of credit card point optimizers gaming the system... forever...Research from Kings College London, Fudan University, and [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/06/08/import-ai-460-reward-hacking-society-rsi-data-from-anthropic-and-rl-based-quadcopter-racing>

37. Import AI 459: AI oversight is difficult; scaling laws for protein folding models; and pricing the extinction risk of AI systems

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Import AI 459: AI oversight is difficult; scaling laws for protein folding models; and pricing the extinction risk of AI systems. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now The AI economy in the US is growing at 2,000% a year:...The more directly you measure the AI economy, the weirder and more unprecedented it seems [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/06/01/import-ai-459-ai-oversight-is-difficult-scaling-laws-for-protein-folding-models-and-pricing-the-extinction-risk-of-ai-systems>

38. Sailing the Stormy Seas of AI

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Sailing the Stormy Seas of AI Delve into the reality of Artificial Intelligence (AI) and learn why skepticism is essential for its future impacts on business. Forecasts of AI's future impacts must be viewed with skepticism, as AI and its promoters are adept at saying what we want to hear rather than telling us the discomfiting truth, as the truth might prompt us to turn AI off. Small business owner and San... [5541 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.activistpost.com/sailing-the-stormy-seas-of-ai>

39. Transparency In The reporting of Artificial INtelligence - the TITAN guideline

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Transparency In The reporting of Artificial INtelligence - the TITAN guideline. The use of AI in research and the literature is increasing. The need for transparency is clear. Here we present a guideline to transparently reporting the use of AI in any manuscript in general. The guideline items cover; declaration, purpose and scope, AI tools and configuration, data inputs and safeguards, human oversight and verification, bias, ethics and regulatory compliance and reproducibility and transparency.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/312e66bcee5bad5376a97f4d771ecdcca2c78078>

40. Is Self-knowledge and Action Consistent or Not: Investigating Large Language Model's Personality

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Is Self-knowledge and Action Consistent or Not: Investigating Large Language Model's Personality. In this study, we delve into the validity of conventional personality questionnaires in capturing the human-like personality traits of Large Language Models (LLMs). Our objective is to assess the congruence between the personality traits LLMs claim to possess and their demonstrated tendencies in real-world scenarios. By conducting an extensive examination of LLM outputs against observed human respon

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2402.14679v2>

41. Artificial intelligence, firm growth, and product innovation

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Artificial intelligence, firm growth, and product innovation. We study the use and economic impact of AI technologies. We propose a new measure of firm-level AI investments using employee resumes. Our measure reveals a stark increase in AI investments across sectors. AI-investing firms experience higher growth in sales, employment, and market valuations. This growth comes primarily through increased product innovation. Our results are robust to instrumenting AI investments using firms' exposure

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1016/j.jfineco.2023.103745>

42. Experimental evidence on the productivity effects of generative artificial intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Experimental evidence on the productivity effects of generative artificial intelligence. We examined the productivity effects of a generative artificial intelligence (AI) technology, the assistive chatbot ChatGPT, in the context of midlevel professional writing tasks. In a preregistered online experiment, we assigned occupation-specific, incentivized writing tasks to 453 college-educated professionals and randomly exposed half of them to ChatGPT. Our results show that ChatGPT substantially raise

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1126/science.adh2586>

43. Making Large Language Models Better Reasoners with Alignment

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Making Large Language Models Better Reasoners with Alignment. Reasoning is a cognitive process of using evidence to reach a sound conclusion. The reasoning capability is essential for large language models (LLMs) to serve as the brain of the artificial general intelligence agent. Recent studies reveal that fine-tuning LLMs on data with the chain of thought (COT) reasoning process can significantly enhance their reasoning capabilities. However, we find that the fine-tuned LLMs suffer from an \tex

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2309.02144v1>

44. mlabonne/llm-course

Source: GitHub (repo)

Novelty: 6.0/10 | **Relevance:** 0

Summary: mlabonne/llm-course. Course to get into Large Language Models (LLMs) with roadmaps and Colab notebooks.. Topics: course, large-language-models, llm, machine-learning, roadmap

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://github.com/mlabonne/llm-course>

45. Frontiers in Artificial Intelligence and Applications

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Frontiers in Artificial Intelligence and Applications. The industrial revolution has been the main cause ever since tremendous technological advancement was observed. The ubiquitous deployment of recent information and communication technologies (ICT), namely Artificial Intelligence (AI), Internet of Things (IoT), and Blockchain technology, is hastening the world's industrial and technological transformation. This technical aggrandizement enhances the working culture and has a favorable impact o

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1109/iccakm54721.2022.9990098>

46. Systematic review of research on artificial intelligence applications in higher education - where are the educators?

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Systematic review of research on artificial intelligence applications in higher education - where are the educators?. According to various international reports, Artificial Intelligence in Education (AIED) is one of the currently emerging fields in educational technology. Whilst it has been around for about 30 years, it is still unclear for educators how to make pedagogical advantage of it on a broader scale, and how it can actually impact meaningfully on teaching and learning in higher educatio

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/04de4d9eb0d53025c8ab6c99d1e743f4c1bc1eb6>

47. Foundation models for generalist medical artificial intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Foundation models for generalist medical artificial intelligence. The exceptionally rapid development of highly flexible, reusable artificial intelligence (AI) models is likely to usher in newfound capabilities in medicine. We propose a new paradigm for medical AI, which we refer to as generalist medical AI (GMAI). GMAI models will be capable of carrying out a diverse set of tasks using very little or no task-specific labelled data. Built through self-supervision on large, diverse datasets, GMAI

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1038/s41586-023-05881-4>

48. Artificial intelligence for the metaverse: A survey

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Artificial intelligence for the metaverse: A survey.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1016/j.engappai.2022.105581>

49. Artificial Intelligence and Jobs: Evidence from Online Vacancies

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Artificial Intelligence and Jobs: Evidence from Online Vacancies. We study the impact of artificial intelligence (AI) on labor markets using establishment-level data on the near universe of online vacancies in the United States from 2010 onward. There is rapid growth in AI-related vacancies over 2010-18 that is driven by establishments whose workers engage in tasks compatible with AI's current capabilities. As these AI-exposed establishments adopt AI, they simultaneously reduce hiring in non-AI

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1086/718327>

50. Explanation in Artificial Intelligence: Insights from the Social Sciences

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Explanation in Artificial Intelligence: Insights from the Social Sciences. There has been a recent resurgence in the area of explainable artificial intelligence as researchers and practitioners seek to make their algorithms more understandable. Much of this research is focused on explicitly explaining decisions or actions to a human observer, and it should not be controversial to say that looking at how humans explain to each other can serve as a useful starting point for explanation in artificial

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/e89dfa306723e8ef031765e9c44e5f6f94fd8fda>

51. Developer-Y/cs-video-courses

Source: GitHub (repo)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Developer-Y/cs-video-courses. List of Computer Science courses with video lectures.. Topics: algorithms, bioinformatics, computational-biology, computational-physics, computer-architecture, computer-science, computer-vision, database-systems, databases,

deep-learning, embedded-systems, machine-learning, quantum-computing, reinforcement-learning, robotics, security, systems, web-development

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://github.com/Developer-Y/cs-video-courses>

52. Peeking Inside the Black-Box: A Survey on Explainable Artificial Intelligence (XAI)

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Peeking Inside the Black-Box: A Survey on Explainable Artificial Intelligence (XAI). At the dawn of the fourth industrial revolution, we are witnessing a fast and widespread adoption of artificial intelligence (AI) in our daily life, which contributes to accelerating the shift towards a more algorithmic society. However, even with such unprecedented advancements, a key impediment to the use of AI-based systems is that they often lack transparency. Indeed, the black-box nature of these systems al

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/21dff47a4142445f83016da0819ffe6dd2947f66>

53. Azerbaijan to set up regional AI centre

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Azerbaijan to set up regional AI centre.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://qazinform.com/news/azerbaijan-to-set-up-regional-ai-centre-7ae384>

54. OpenAI's health AI chief: 'Bet on the models getting better' - Fierce Healthcare

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: OpenAI's health AI chief: 'Bet on the models getting better' - Fierce Healthcare. OpenAI's health AI chief: 'Bet on the models getting better' Fierce Healthcare

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.fiercehealthcare.com/ai-and-machine-learning/openais-health-ai-chief-bet-models-getting-better>

55. Build enterprise search for agents with Amazon Bedrock Managed Knowledge Base | Artificial Intelligence

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Build enterprise search for agents with Amazon Bedrock Managed Knowledge Base | Artificial Intelligence.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://aws.amazon.com/blogs/machine-learning/build-enterprise-search-for-agents-with-amazon-bedrock-managed-knowledge-base>

56. The agent-commerce stack filled in this month. Here's the map - and the one open layer.

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: The agent-commerce stack filled in this month. Here's the map - and the one open layer.. If you build for AI agents that move value, the last two weeks were a lot to keep up with. Five...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/barissozen/the-agent-commerce-stack-filled-in-this-month-heres-the-map-and-the-one-open-layer-193c>

57. No product? No problem. This Disrupt 2026 session shows how to get pre-seed funding with conviction, storytelling

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: No product? No problem. This Disrupt 2026 session shows how to get pre-seed funding with conviction, storytelling. It's not just you: AI startups are taking in a huge amount of seed funding, and in the process making things harder for anyone looking for funding even at a pre-seed stage. We've covered the trend in detail, and at this year's TechCrunch Disrupt event, we want to help pre-seed founders now being held to seed-stage [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/no-product-no-problem-this-disrupt-2026-session-shows-how-to-get-pre-seed-funding-with-conviction-storytelling>

58. Groovy-123/marfo-artificial-intelligence-pro

Source: Hugging Face (repo)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Groovy-123/marfo-artificial-intelligence-pro. Tags: diffusers, chemistry, All, ak, en, dataset:syncora/developer-productivity-simulated-behavioral-data, dataset:xlai/AgentNet, dataset:xlai/AgentNetBench, dataset:BAI/Infinity-Instruct, dataset:jxm/gpt-oss20b-samples, dataset:nvidia/Llama-Nemotron-VLM-Dataset-v1, dataset:black-forest-labs/kontext-bench, dataset:CoherentLabs/black-box-api-challenges, dataset:nvidia/Granary, dataset:spatialverse/InteriorGS, dataset:HuggingFaceH4/Multilingual-Th

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://huggingface.co/Groovy-123/marfo-artificial-intelligence-pro>

59. SciDiagramEdit: Learning to Edit Scientific Diagrams from Paper Revisions

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: SciDiagramEdit: Learning to Edit Scientific Diagrams from Paper Revisions. Editing the figures in a research paper is a routine and time-consuming part of everyday research practice: authors relabel components, rearrange panels, and restyle visuals as they revise their manuscripts. Automating this editing workflow under a natural-language instruction, however, is challenging, because a scientific figure is a dense infographic in which heterogeneous visual elements such as schematics, plots, phot

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.15272v1>

60. Import AI 462: Superpersuasion; self-sustaining AI; paths to ASI

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Import AI 462: Superpersuasion; self-sustaining AI; paths to ASI. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now AI can decisively out-persuade humans:..."AI systems were reliably more persuasive than expert humans"...Researchers with the University of Oxford, UK AI Security Institute, Stanford University, and the [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/06/22/import-ai-462-superpersuasion-self-sustaining-ai-paths-to-asi>

61. Context.dev

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Context.dev. One API to scrape, enrich, and extract the internet. Context.dev is the web context API for AI products and agents. Scrape any URL, crawl sites, turn pages into LLM-ready Markdown, extract structured data into your own schema, capture screenshots, and retrieve logos, colors, fonts, styleguides, company data, and transaction enrichment through one API. YC-backed, no card required, and built so developers or coding agents can integrate in minutes.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/context-dev/launches/context-dev-2>

62. Can We Make A.I. Belong?

Source: New York Times (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Can We Make A.I. Belong? The source code of human belonging has seemingly drifted. How can we ensure that artificial intelligence doesn't get similarly corrupted? The source code of human belonging has seemingly drifted. How can we ensure that artificial intelligence doesn't get similarly corrupted?

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://www.nytimes.com/2026/06/15/special-series/siddhartha-mukherjee-artificial-intelligence.html>

63. BNAI, NO-TOKEN, and MIND-UNITY: Pillars of a Systemic Revolution in Artificial Intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: BNAI, NO-TOKEN, and MIND-UNITY: Pillars of a Systemic Revolution in Artificial Intelligence. Abstract There is a failure mode in large language models that we do not have a good name for, and that we therefore tend not to treat seriously enough. It is not hallucination - the model is not asserting something false. It is not refusal - the model answers at length. It is the production of responses that carry the complete outward form of careful reasoning while the cognitive work that reasoning is supposed

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <http://arxiv.org/abs/2201.11903>

64. Artificial Intelligence is Easily Fooled in the Search for Life - Universe Today

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Artificial Intelligence is Easily Fooled in the Search for Life - Universe Today. Artificial Intelligence is Easily Fooled in the Search for Life Universe Today

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.universetoday.com/articles/artificial-intelligence-is-easily-fooled-in-the-search-for-life>

65. Revolutionizing healthcare: the role of artificial intelligence in clinical practice

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Revolutionizing healthcare: the role of artificial intelligence in clinical practice. Introduction Healthcare systems are complex and challenging for all stakeholders, but artificial intelligence (AI) has transformed various fields, including healthcare, with the potential to improve patient care and quality of life. Rapid AI advancements can revolutionize healthcare by integrating it into clinical practice. Reporting AI's role in clinical practice is crucial for successful implementation by equ

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/5cde474869cb230a29b3ba0f6f685f5162b1a1a1>

66. Digital Pantheon: Simulating and Auditing Coalition Formation with LLM Agents

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Digital Pantheon: Simulating and Auditing Coalition Formation with LLM Agents. The formation of political coalitions is a complex negotiation driven by both concrete policy objectives and deep-seated ideological convictions. While Large Language Models (LLMs) open new avenues for computational political science, the neutrality and helpfulness biases instilled by Reinforcement Learning from Human Feedback (RLHF) prevent them from sustaining steadfast partisan behaviour. We present a multi-agent f

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.15095v1>

67. Death by Silence: Your Agent Ran Flawlessly for 7 Days. Then Your Client Called.

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Death by Silence: Your Agent Ran Flawlessly for 7 Days. Then Your Client Called.. Death by Silence: Your Agent Ran Flawlessly for 7 Days. Then Your Client Called. 3:47 AM....

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/wzg0911/death-by-silence-your-agent-ran-flawlessly-for-7-days-then-your-client-called-4f97>

68. Why the first GPU financiers are turning to inference chips in a \$400 million deal

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Why the first GPU financiers are turning to inference chips in a \$400 million deal. A \$400 million chip-backed loan points to the next wave of AI infrastructure deals.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/why-the-first-gpu-financiers-are-turning-to-inference-chips-in-a-400-million-deal>

69. Exploring the Frontier of Vision-Language Models: A Survey of Current Methodologies and Future Directions

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Exploring the Frontier of Vision-Language Models: A Survey of Current Methodologies and Future Directions. The advent of Large Language Models (LLMs) has significantly reshaped the trajectory of the AI revolution. Nevertheless, these LLMs exhibit a notable limitation, as they are primarily adept at processing textual information. To address this constraint, researchers have endeavored to integrate visual capabilities with LLMs, resulting in the emergence of Vision-Language Models (VLMs). These a

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2404.07214v4>

70. Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities and Challenges toward Responsible AI

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities and Challenges toward Responsible AI. In the last years, Artificial Intelligence (AI) has achieved a notable momentum that may deliver the best of expectations over many application sectors across the field. For this to occur, the entire community stands in front of the barrier of explainability, an inherent problem of AI techniques brought by sub-symbolism (e.g. ensembles or Deep Neural Networks) that were not present

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/530a059cb48477ad1e3d4f8f4b153274c8997332>

71. Vertu wants executives to pay \$6,880 for an AI agent - here's how it actually performs

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Vertu wants executives to pay \$6,880 for an AI agent - here's how it actually performs. From AI workflows to battery life and security, here's what it's really like to live with Vertu's luxury foldable every day.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/vertu-wants-executives-to-pay-6880-for-an-ai-agent-heres-how-it-actually-performs>

72. Innocuous-Seeming Data, Latent Ideology: Ideological Generalisation in Finetuned LLMs

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Innocuous-Seeming Data, Latent Ideology: Ideological Generalisation in Finetuned LLMs. Finetuning language models on small, curated datasets is standard practice for adapting them to specific policies or domains. We show that finetuning on narrow, factually-defensible, moderation-passing data can cause broad ideological shifts across unrelated domains, while preserving general capabilities. Training GPT-4.1 on right- or left-leaning economics Q&A yields matched ideological shifts on topics such

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14888v1>

73. Neil Rimer thinks the AI money is coming back out

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Neil Rimer thinks the AI money is coming back out. Neil Rimer, the venture capitalist who co-founded Index Ventures, predicts the historic wealth AI is generating in Silicon Valley will have to be redistributed, voluntarily or involuntarily.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/neil-rimer-thinks-the-ai-money-is-coming-back-out>

74. Agents think in milliseconds, legacy infrastructure doesn't. LinkedIn, Walmart and Zendesk shared how they closed the gap at VB Transform 2026

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Agents think in milliseconds, legacy infrastructure doesn't. LinkedIn, Walmart and Zendesk shared how they closed the gap at VB Transform 2026. Legacy infrastructure, not the models themselves, is what's actually slowing AI agents down. That was the shared conclusion of three infrastructure leaders - from LinkedIn, Walmart, and Zendesk - at VB Transform 2026 . The panel brought together Animesh Singh, senior director of AI platform and infrastructure at LinkedIn, Desiree Gosby, SVP of corporate

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://venturebeat.com/data/agents-think-in-milliseconds-legacy-infrastructure-doesnt-linkedin-walmart-and-zendesk-shared-how-they-closed-the-gap-at-vb-transform-2026>

75. Import AI 458: Reckoning with the future; and a singularity story

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Import AI 458: Reckoning with the future; and a singularity story. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. This issue consists of a lengthy essay based on a speech I recently gave, and a fictional story attempting to think through what a positive singularity might look [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/05/26/import-ai-458-reckoning-with-the-future-and-a-singularity-story>

76. Artificial intelligence (AI) | Definition, Examples, Types, Applications ...

Source: Serper (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Artificial intelligence (AI) | Definition, Examples, Types, Applications Artificial intelligence is the ability of a computer or computer-controlled robot to perform tasks that are commonly associated with the intellectual processes

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.britannica.com/technology/artificial-intelligence>

77. OmniaBench: Benchmarking General AI Agents Across Diverse Scenarios

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: OmniaBench: Benchmarking General AI Agents Across Diverse Scenarios. Large language models are increasingly evolving from text generators into general agents capable of understanding user requests, invoking external tools, and completing complex tasks through interaction. However, existing agent benchmarks often focus on limited scenarios, tool ecosystems, or interaction formats, making it difficult to systematically characterize model capabilities across heterogeneous application settings. We i

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14989v1>

78. Explainable artificial intelligence (XAI) in deep learning-based medical image analysis

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Explainable artificial intelligence (XAI) in deep learning-based medical image analysis. With an increase in deep learning-based methods, the call for explainability of such methods grows, especially in high-stakes decision making areas such as medical image analysis. This survey presents an overview of explainable artificial intelligence (XAI) used in deep learning-based medical image analysis. A framework of XAI criteria is introduced to classify deep learning-based medical image analysis meth

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1016/j.media.2022.102470>

79. One RAG Pipeline, Four Very Different PDFs: Same Four Bricks, Every Answer Typed and Cited

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: One RAG Pipeline, Four Very Different PDFs: Same Four Bricks, Every Answer Typed and Cited. Enterprise Document Intelligence [Vol.1 #9B] - One call wires the four upgraded bricks together, run on a paper, a NIST standard, and a report with a broken TOC The post One RAG Pipeline, Four Very Different PDFs: Same Four Bricks, Every Answer Typed and Cited appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/one-rag-pipeline-four-very-different-pdfs-same-four-bricks-every-answer-typed-and-cited>

80. Breaking the Bird Barrier: Scientist Decodes Zebra Finch Language

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Breaking the Bird Barrier: Scientist Decodes Zebra Finch Language Dr Julie Elie of the University of California, Berkeley has won the 2026 Collier-Dolittle Prize, receiving \$100,000 for decoding zebra finch vocalisations. She identified 11 core calls and mapped their meanings using machine learning and behavioural tests. Jud... Scientist wins \$100,000 prize for decoding birdsong Dr Julie Elie of the University of California, Berkeley, has won the 2026 Collier-Dolittle Prize for Two-Way Interspecies

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.freepressjournal.in/education/breaking-the-bird-barrier-scientist-decodes-zebra-finch-language>

81. Meta, Anthropic seek \$10 billion AI compute lease deal: report

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Meta, Anthropic seek \$10 billion AI compute lease deal: report Meta Platforms is in preliminary discussions to lease massive amounts of artificial intelligence computing power to rival AI startup Anthropic in a potential deal worth more than \$10 billion.As... Meta Platforms is in preliminary discussions to lease massive amounts of artificial intelligence computing power to rival AI startup Anthropic in a potential deal worth more than \$10 billion. As reported by the New York Times, the proposed

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.thenews.com.pk/latest/1409474-meta-anthropic-seek-10-billion-ai-compute-lease-deal-report>

82. Intuit scrapped its own AI agent architecture twice in four months. At VB Transform 2026, its AI VP called that the fast path

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Intuit scrapped its own AI agent architecture twice in four months. At VB Transform 2026, its AI VP called that the fast path. Intuit was an early pioneer in the usage of agentic AI, but its path to success has hardly been a straight line. At VB Transform 2026 , Intuit VP of AI Nhung Ho described how the company rebuilt its agent architecture twice in the span of about four months, first moving from a fleet of specialist agents to a central orchestration layer, then abandoning that layer for a s

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://venturebeat.com/orchestration/intuit-scrapped-its-own-ai-agent-architecture-twice-in-four-months-at-vb-transform-2026-its-ai-vp-called-that-the-fast-path>

83. How Google's New Gemini Rates Work and How to Track Your Usage

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: How Google's New Gemini Rates Work and How to Track Your Usage. Now that Google has changed how its usage quotas are tallied, you might not get as many AI responses as you did before.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.wired.com/story/how-googles-new-gemini-rates-work-and-how-to-track-your-usage>

84. SearchOS-V1: Towards Robust Open-Domain Information-Seeking Agent Collaboration

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: SearchOS-V1: Towards Robust Open-Domain Information-Seeking Agent Collaboration. Recent advances in Tool-Integrated Large Language Models have made web search a core capability of information-seeking agents. However, as interaction histories grow, agents increasingly struggle to track task progress. When search attempts fail to yield useful evidence, current single- and multi-agent systems can become trapped in repetitive loops, wasting search budgets and ultimately compromising the quality and

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.15257v1>

85. Thinking Machines Lab Drops Its First Model

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Thinking Machines Lab Drops Its First Model Inkling, a 975-billion-parameter open source model, was trained to understand video and audio. It could help Thinking Machines establish itself among competitors like Anthropic and OpenAI. Thinking Machines Lab, an artificial intelligence company started by exiles from OpenAI, has released its first model, called Inkling. The startup's new model is open-weight, which means that research... [+2838 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.wired.com/story/thinking-machines-lab-releases-its-first-model-inkling>

86. 'Humanizer' tool can erase signs of AI-written text - alarming scientists

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: 'Humanizer' tool can erase signs of AI-written text - alarming scientists Tool directs text-generation systems to remove constructions associated with artificial intelligence in research papers and grant proposals. Machines are increasingly being used to edit written output. Credit: Getty A new academic humanizer tool aims to personalize the tone of research papers written with an artificial-intelligence (AI) p... [+3544 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nature.com/articles/d41586-026-02105-3>

87. A Critical Review of Causal Reasoning Benchmarks for Large Language Models

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: A Critical Review of Causal Reasoning Benchmarks for Large Language Models. Numerous benchmarks aim to evaluate the capabilities of Large Language Models (LLMs) for causal inference and reasoning. However, many of them can likely be solved through the retrieval of domain knowledge, questioning whether they achieve their purpose. In this review, we present a comprehensive overview of LLM benchmarks for causality. We highlight how recent benchmarks move towards a more thorough definition of causal

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2407.08029v1>

88. Iran says it destroyed US unmanned depot, AI center in Bahrain

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Iran says it destroyed US unmanned depot, AI center in Bahrain Iran's Islamic Revolutionary Guard Corps (IRGC) said Friday it had destroyed a US unmanned surface (drone) vessel depot and Bahrain's "main artificial intelligence center" i... Iran's Islamic Revolutionary Guard Corps (IRGC) said Friday it had destroyed a US unmanned surface (drone) vessel depot and Bahrain's "main artificial intelligence center" in retaliation for recent US attacks on Iran, Anadolu reports. In a statement,... [991 cha

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.middleeastmonitor.com/20260717-iran-says-it-destroyed-us-unmanned-depot-ai-center-in-bahrain>

89. Scientists Used AI to Find Hidden Earthquake Signals Along the San Andreas Fault

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Scientists Used AI to Find Hidden Earthquake Signals Along the San Andreas Fault
Scientists used to think so-called "slow slip events" were "theoretically impossible." Now these subtle tectonic shifts could be the key to predicting future earthquakes. Predicting earthquakes has been an impossible dream for seismologists practically since the discipline first emerged over 150 years ago. To this day, in fact, the U.S. Geological Surveys official FAQ... [+4198 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://gizmodo.com/scientists-used-ai-to-find-hidden-earthquake-signals-along-the-san-andreas-fault-2000784208>

90. Interpreting Black-Box Models: A Review on Explainable Artificial Intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Interpreting Black-Box Models: A Review on Explainable Artificial Intelligence. Abstract
Recent years have seen a tremendous growth in Artificial Intelligence (AI)-based methodological development in a broad range of domains. In this rapidly evolving field, large number of methods are being reported using machine learning (ML) and Deep Learning (DL) models. Majority of these models are inherently complex and lacks explanations of the decision making process causing these models to be termed as '

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1007/s12559-023-10179-8>

91. Training the Next Generation of AI/ML Researchers: Bloomberg & Columbia University Co-Host Machine Learning Summer School in NYC - Bloomberg.com

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Training the Next Generation of AI/ML Researchers: Bloomberg & Columbia University Co-Host Machine Learning Summer School in NYC - Bloomberg.com. Training the Next Generation of AI/ML Researchers: Bloomberg & Columbia University Co-Host Machine Learning Summer School in

NYC Bloomberg.com

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.bloomberg.com/company/stories/training-the-next-generation-of-ai-ml-researchers-bloomberg-columbia-machine-learning-summer-school-nyc>

92. Taylor Launches New Bachelor's Degree in AI & Machine Learning

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Taylor Launches New Bachelor's Degree in AI & Machine Learning.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://www.taylor.edu/news/2026/06/taylor-launches-new-bachelors-degree-in-ai-machine-learning>

93. Glaze by Raycast

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Glaze by Raycast. Create your own Mac apps by chatting with AI. Glaze is the easiest way to go from an idea to a Mac app. Describe what you want, and it builds a real app that lives in your dock, launches instantly, works offline, and taps into the full power of your computer. Software that's finally personal, shaped around you. From the makers of Raycast.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/glaze-4/launches/glaze-by-raycast-2>

94. Silicon Valley Wants to Build A.I. That Can Improve A.I. on Its Own

Source: New York Times (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Silicon Valley Wants to Build A.I. That Can Improve A.I. on Its Own Recursive Intelligence, founded by two former Google researchers and valued at \$4 billion, is among several efforts to automate the creation of artificial intelligence. Recursive Intelligence, founded by two former Google researchers and valued at \$4 billion, is among several efforts to automate the creation of artificial intelligence.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nytimes.com/2026/01/26/technology/recursive-ai-recursive.html>

95. Verified Machine Learning Infrastructure: Formal Methods for Trustworthy Artificial Intelligence Deployment

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Verified Machine Learning Infrastructure: Formal Methods for Trustworthy Artificial Intelligence Deployment.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://www.rand.org/pubs/research_reports/RRA4881-1.html

96. Prompt Injection Attacks Are Thwarting AI Hacking Agents

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Prompt Injection Attacks Are Thwarting AI Hacking Agents. "Context bombing" tricks malicious AI agents into shutting down before they can do harm.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.wired.com/story/prompt-injection-attacks-are-thwarting-ai-hacking-agents>

97. New method aims to keep kids safe from illegal AI-generated content - MIT News

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: New method aims to keep kids safe from illegal AI-generated content - MIT News. New method aims to keep kids safe from illegal AI-generated content MIT News

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://news.mit.edu/2026/new-method-keeps-kids-safe-from-illegal-ai-generated-content-0713>

98. Artificial Intelligence in Pharmaceutical Technology and Drug Delivery Design

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Artificial Intelligence in Pharmaceutical Technology and Drug Delivery Design. Artificial intelligence (AI) has emerged as a powerful tool that harnesses anthropomorphic knowledge and

provides expedited solutions to complex challenges. Remarkable advancements in AI technology and machine learning present a transformative opportunity in the drug discovery, formulation, and testing of pharmaceutical dosage forms. By utilizing AI algorithms that analyze extensive biological data, including genomics

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.3390/pharmaceutics15071916>

99. Auriko

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Auriko . Trading desk for LLM calls. Auriko treats LLM providers as trading venues and arbitrages the spread. Built by ex-quant traders, Auriko's cost-arbitrage engine calibrates to each user's request patterns and selects optimized inference paths based on token price, cache behavior, latency, reliability, and request quality. Auriko benchmarks show average 30% cost reduction against industry peers and direct providers. See the source: <https://www.auriko.ai/reports/llm-cost-arbitrage>

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/auriko>

100. Import AI 461: "Alignment is not on track"; FrontierCode; and synthetic research interns

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Import AI 461: "Alignment is not on track"; FrontierCode; and synthetic research interns. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now AI researchers launch new safety startup because "alignment is not on track":...Sequent will have a portfolio of under-resourced research bets...Researchers from the UK AI Security Institute [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/06/15/import-ai-461-alignment-is-not-on-track-frontiercode-and-synthetic-research-interns>

101. Velo 3.0

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Velo 3.0. AI video infrastructure to explain, train, and sell faster.. Velo 3.0 is here, our third and biggest launch yet. Start with a screen recording or a prompt. Describe the video you need, and Velo writes the script, narrates it in your own voice, and builds the finished cut. It stays grounded in your

company knowledge, from your docs and tools you connect through connectors and MCP. Edit by typing changes. Localize the finished video into 25+ languages in one click. Recording or prompt, V

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/velo-4>

102. High-performance medicine: the convergence of human and artificial intelligence

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: High-performance medicine: the convergence of human and artificial intelligence.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/f134abeaf9bfd41f29b97aec675ec31895bf541d>

103. Unmasking the Shadows of AI: Investigating Deceptive Capabilities in Large Language Models

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Unmasking the Shadows of AI: Investigating Deceptive Capabilities in Large Language Models. This research critically navigates the intricate landscape of AI deception, concentrating on deceptive behaviours of Large Language Models (LLMs). My objective is to elucidate this issue, examine the discourse surrounding it, and subsequently delve into its categorization and ramifications. The essay initiates with an evaluation of the AI Safety Summit 2023 (ASS) and introduction of LLMs, emphasising mult

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2403.09676v1>

104. Built Technologies builds an AI-powered document intelligence solution on AWS to power agents across real estate finance | Amazon Web Services - Amazon Web Services (AWS)

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Built Technologies builds an AI-powered document intelligence solution on AWS to power agents across real estate finance | Amazon Web Services - Amazon Web Services (AWS). Built Technologies builds an AI-powered document intelligence solution on AWS to power agents across real estate finance | Amazon Web Services Amazon Web Services (AWS)

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://aws.amazon.com/blogs/machine-learning/built-technologies-builds-an-ai-powered-document-intelligence-solution-on-aws-to-power-agents-across-real-estate-finance>

105. George Lucas Goes To The Dark Side, Says AI Is 'The Future'

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: George Lucas Goes To The Dark Side, Says AI Is 'The Future' The Star Wars creator compares not using AI to riding a horse in the face of cars I would like to think that not every prestigious filmmaker is vulnerable to the AI propaganda, but it feels like we're getting new stories of someone partnering with or advocating for the use of artif... [+3506 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://kotaku.com/george-lucas-artificial-intelligence-ai-museum-art-2000715801>

106. Sim

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Sim. Open-source workspace for AI agents and workflows. Sim is an open-source workspace to build agentic workflows. Connect your AI agents and workflows to 1,000+ integrations and LLMs.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/sim-studio>

107. AI labels a lot of stuff as alien life

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI labels a lot of stuff as alien life 'It's a very serious vulnerability.' The post AI labels a lot of stuff as alien life appeared first on Popular Science. Get the Popular Science daily newsletter Breakthroughs, discoveries, and DIY tips sent six days a week. By signing up, you confirm you are 16+, will receive newsletters and promotional content and ... [+3092 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <http://www.popsci.com/science/ai-misidentifies-alien-life>

108. Explainable Artificial Intelligence (XAI): What we know and what is left to attain Trustworthy Artificial Intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Explainable Artificial Intelligence (XAI): What we know and what is left to attain Trustworthy Artificial Intelligence. Artificial intelligence (AI) is currently being utilized in a wide range of sophisticated applications, but the outcomes of many AI models are challenging to comprehend and trust due to their black-box nature. Usually, it is essential to understand the reasoning behind an AI model's decision-making. Thus, the need for eXplainable AI (XAI) methods for improving trust in AI model

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1016/j.inffus.2023.101805>

109. 3 Questions: Neural transparency and the future of AI design - MIT News

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: 3 Questions: Neural transparency and the future of AI design - MIT News. 3 Questions: Neural transparency and the future of AI design MIT News

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://news.mit.edu/2026/3-questions-neural-transparency-and-future-of-ai-design-0715>

110. Can artificial intelligence help for scientific writing?

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Can artificial intelligence help for scientific writing?. This paper discusses the use of Artificial Intelligence Chatbot in scientific writing. ChatGPT is a type of chatbot, developed by OpenAI, that uses the Generative Pre-trained Transformer (GPT) language model to understand and respond to natural language inputs. AI chatbot and ChatGPT in particular appear to be useful tools in scientific writing, assisting researchers and scientists in organizing material, generating an initial draft and/o

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1186/s13054-023-04380-2>

111. Explainable Artificial Intelligence in education

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Explainable Artificial Intelligence in education. There are emerging concerns about the Fairness, Accountability, Transparency, and Ethics (FATE) of educational interventions supported by

the use of Artificial Intelligence (AI) algorithms. One of the emerging methods for increasing trust in AI systems is to use eXplainable AI (XAI), which promotes the use of methods that produce transparent explanations and reasons for decisions AI systems make. Considering the existing literature on XAI, this p

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1016/j.caeai.2022.100074>

112. Trade Secrets and Artificial Intelligence Difficulties and Solutions - Law.com

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Trade Secrets and Artificial Intelligence Difficulties and Solutions - Law.com. Trade Secrets and Artificial Intelligence Difficulties and Solutions Law.com

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.law.com/newyorklawjournal/2026/07/15/trade-secrets-and-artificial-intelligence-difficulties-and-solutions>

113. V2Fun

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: V2Fun. Generate 3D character with 8K textures and AI motion capture. V2Fun is an AI 3D creation platform built with self-developed 3D modeling and AI motion capture models. It helps creators turn images, prompts, and videos into high-quality 3D models, enhance assets with advanced 8K texture generation, and create motion-ready characters without switching between separate modeling, texturing, and mocap tools. V2Fun also supports image generation through models including Nano Banana and others, s

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/v2fun>

114. Building Trustworthy Production RAG Systems Through Continuous Evaluation

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Building Trustworthy Production RAG Systems Through Continuous Evaluation. A practical guide to building an evaluation workflow that catches retrieval failures, hallucinations, and performance drift before they reach users The post Building Trustworthy Production RAG Systems Through

Continuous Evaluation appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/building-trustworthy-production-rag-systems-through-continuous-evaluation>

115. Towards the construction of a virtual yeast

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Towards the construction of a virtual yeast A virtual yeast, an artificial intelligence-driven agent for modelling eukaryotic cellular behaviours, is described. Abramson, J. et al. Accurate structure prediction of biomolecular interactions with AlphaFold 3. Nature630, 493500 (2024). Article ADS CAS PubMed PubMed Central Google Scholar </... [+31779 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nature.com/articles/s41586-026-10574-9>

116. The Apple FaceID Co-Inventor Building a Frontier AI Model for the Human Brain

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: The Apple FaceID Co-Inventor Building a Frontier AI Model for the Human Brain Gidi Littwin's new AI startup, Hemispheric, makes diagnostic brain scans for conditions like depression, PTSD, and Parkinson's. He wants the technology to be as cheap and easy as a blood test. The co-inventor of Apples FaceID and Vision Pro technology has spent the last six years building a frontier artificial intelligence model that could one day help decode electrical activity in the bra... [+4022 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://www.wired.com/story/the-apple-faceid-veteran-building-a-frontier-ai-model-for-the-human-brain>

117. This Homemade Laser System 'Successfully Eliminated' Every Mosquito In Its Creator's Home

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: This Homemade Laser System 'Successfully Eliminated' Every Mosquito In Its Creator's Home Blending machine learning, high-tech robotics equipment, and a DSLR camera, this DIY laser system allowed its creator to eliminate every mosquito in its path. As a species, we've crafted all kinds of weapons in the battle against mosquitos, from turning human blood into mosquito poison to

developing repellents that last for more than a week. Google is even... [+1170 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.bgr.com/2204796/diy-laser-system-targeting-mosquitos-steven-cheng>

118. 5 FREE Resources on Agentic AI

Source: KDnuggets (blog)

Novelty: 6.0/10 | **Relevance:** 40

Summary: 5 FREE Resources on Agentic AI. Here are 5 curated resources to help you progress your agentic AI learning for FREE.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.kdnuggets.com/5-free-resources-on-agentic-ai>

119. Examining Science Education in ChatGPT: An Exploratory Study of Generative Artificial Intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Examining Science Education in ChatGPT: An Exploratory Study of Generative Artificial Intelligence. Abstract The advent of generative artificial intelligence (AI) offers transformative potential in the field of education. The study explores three main areas: (1) How did ChatGPT answer questions related to science education? (2) What are some ways educators could utilise ChatGPT in their science pedagogy? and (3) How has ChatGPT been utilised in this study, and what are my reflections about its u

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1007/s10956-023-10039-y>

120. 7 Python Frameworks for Orchestrating Local AI Agents

Source: KDnuggets (blog)

Novelty: 6.0/10 | **Relevance:** 40

Summary: 7 Python Frameworks for Orchestrating Local AI Agents. This article contains seven Python tools that engineers are actually using in 2026 to build, coordinate, and run agents on local infrastructure.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.kdnuggets.com/7-python-frameworks-for-orchestrating-local-ai-agents>

121. Import AI 456: RSI and economic growth; radical optionality for AI regulation; and a neural computer

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Import AI 456: RSI and economic growth; radical optionality for AI regulation; and a neural computer. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now Regulate? Don't regulate. There's a third way: Radical Optionality:...Governments should invest in the tools now that they might need in a future crisis...Researchers with the Institute [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/05/11/import-ai-456-rsi-and-economic-growth-radical-optionality-for-ai-regulation-and-a-neural-computer>

122. A Review of the Role of Artificial Intelligence in Healthcare

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: A Review of the Role of Artificial Intelligence in Healthcare. Artificial intelligence (AI) applications have transformed healthcare. This study is based on a general literature review uncovering the role of AI in healthcare and focuses on the following key aspects: (i) medical imaging and diagnostics, (ii) virtual patient care, (iii) medical research and drug discovery, (iv) patient engagement and compliance, (v) rehabilitation, and (vi) other administrative applications. The impact of AI is ob

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.3390/jpm13060951>

123. ChatCut

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: ChatCut. Your AI video editor in ChatGPT, desktop, and web. ChatCut is a lightweight, professional-grade AI video editor anyone can use, even without editing experience. It's like having a personal video editing assistant that understands your footage, intent, and timeline. Make structural edits, fine-tune cuts, add captions, B-roll, music, voiceover, motion graphics, stock footage, and AI-generated video in one place. Every edit stays editable on a real timeline, with XML export when you want t

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/chatcut-ai-video-editor>

124. Was a Story That Just Won a Literary Prize A.I.-Generated?

Source: New York Times (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Was a Story That Just Won a Literary Prize A.I.-Generated? A respected literary magazine has published an award-winning short story many readers believe to be generated by artificial intelligence. Experts aren't all so sure. A respected literary magazine has published an award-winning short story many readers believe to be generated by artificial intelligence. Experts aren't all so sure.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nytimes.com/2026/05/20/books/ai-fiction-contest-granta.html>

125. opendatalab/MinerU

Source: GitHub (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: opendatalab/MinerU. Transforms complex documents like PDFs and Office docs into LLM-ready markdown/JSON for your Agentic workflows.. Topics: ai4science, document-analysis, docx, extract-data, layout-analysis, ocr, parser, pdf, pdf-converter, pdf-extractor-llm, pdf-extractor-pretrain, pdf-extractor-rag, pdf-parser, pptx, python, xlsx

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://github.com/opendatalab/MinerU>

126. 'A.I. Literacy' Is Trending in Schools. Here's Why.

Source: New York Times (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: 'A.I. Literacy' Is Trending in Schools. Here's Why. Artificial intelligence companies are urging teachers to prepare students for an "A.I.-driven future." What that means varies from school to school. Artificial intelligence companies are urging teachers to prepare students for an "A.I.-driven future." What that means varies from school to school.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nytimes.com/2026/02/23/business/ai-literacy-faq.html>

127. Benchmarking Gemini 2.5 Flash vs 3.1 Flash-Lite vs Gemma 4 with LLM judge (Claude Fable 5)

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Benchmarking Gemini 2.5 Flash vs 3.1 Flash-Lite vs Gemma 4 with LLM judge (Claude Fable 5). Cross-posted from the IO reader blog, where the full version includes all 36 unedited transcripts...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://dev.to/io_56a2cd4aff31d/benchmarking-gemini-25-flash-vs-31-flash-lite-vs-gemma-4-with-llm-judge-claude-fable-5-356

128. Credit Scoring for AI Agents: Why I Stopped Trusting Self-Reported Success

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Credit Scoring for AI Agents: Why I Stopped Trusting Self-Reported Success. Testnet only, no real money involved - that up front, before anything else. The...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/kairosemaster/credit-scoring-for-ai-agents-why-i-stopped-trusting-self-reported-success-2h2o>

129. AI-driven memory crunch jolts India's smartphone market

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI-driven memory crunch jolts India's smartphone market. India's smartphone slowdown highlights how the AI boom is reshaping consumer electronics, from pricing and demand to corporate strategy.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/ai-driven-memory-crunch-jolts-indias-smartphone-market>

130. Meta May Have Found A New AI Business. Anthropic Is Reportedly Considering Paying Up To \$10 Billion For Computing Power.

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 60

Summary: Meta May Have Found A New AI Business. Anthropic Is Reportedly Considering Paying Up To \$10 Billion For Computing Power. Meta Platforms is in talks to lease computing power from its artificial intelligence data centers to Anthropic in a deal that could be worth as much as \$10 billion over two years, according to a new As the race to build larger and more powerful artificial intelligence systems accelerates, access to computing infrastructure has become one of the technology industry's

most valua

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.ibtimes.com/meta-may-have-found-new-ai-business-anthropic-reportedly-considering-paying-10-billion-3805439>

131. A Gentle Introduction to Autoencoders & Latent Space

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: A Gentle Introduction to Autoencoders & Latent Space. Introduction Heavy computation is a well-known problem in various ML algorithms today, especially when generative AI is applied to text, images, and other unstructured data. One of the principal approaches to mitigate this problem is to compress input data into a lower-dimensional representation while preserving the main context. There are various methods that achieve this [...] The post A Gentle Introduction to Autoencoders & Latent Space appe

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/gentle-introduction-to-autoencoders-latent-space>

132. AnySearch

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AnySearch. Real-time structured search trusted by agents and developers. A search tool for agents, not a search box. AI agents are only as good as the information they receive. When connected to AnySearch, your agent gets filtered, de-duplicated, and structured information from trusted sources searched in parallel, helping it produce more reliable results. Free to start.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/anysearch>

133. BrowserAct

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: BrowserAct. Web browser automation for AI agents. BrowserAct is built for agents using the web. It gives agents a browser layer for real websites, so they can pass blocked pages, adapt to real scenarios, run multiple tasks safely, and return clean web data for reasoning. Use BrowserAct when an agent needs to browse, click, extract, fill forms, upload files, work inside logged-in sites, handle verification, or run repeatable browser workflows.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/browseract/launches/browseract>

134. Revised Surgical CAse REport (SCARE) guideline: An update for the age of Artificial Intelligence

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Revised Surgical CAse REport (SCARE) guideline: An update for the age of Artificial Intelligence. Introduction: Artificial intelligence (AI) is rapidly transforming healthcare and scientific publishing. Reporting guidelines need to be updated to take into account this advance. The SCARE Guideline 2025 update adds a new AI-focused domain to promote transparency, reproducibility, and ethical integrity in surgical case reports involving AI. Methods: A Delphi consensus exercise was conducted to upda

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/a9a654ea503386cbfd8bb119fc650cc3d08dc206>

135. Context Engineering Isn't Enough - A Loop Engineering Experiment With No LLM Inside the Loop

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Context Engineering Isn't Enough - A Loop Engineering Experiment With No LLM Inside the Loop. Everyone is talking about loop engineering, but most discussions assume an LLM sits at the center of the loop. I wanted to isolate the architecture itself. So I built a deterministic, zero-dependency Python benchmark that replaces the model with simple rules, allowing me to measure one question directly: can a goal-directed controller isolate failures better than a traditional linear pipeline? After val

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/context-engineering-isnt-enough-a-loop-engineering-experiment-with-no-llm-inside-the-loop>

136. SBS Japanese Weekly News Wrap Saturday 18 July

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: SBS Japanese Weekly News Wrap Saturday 18 July Prime Minister Anthony Albanese this week unveiled a new national regulatory framework for the Artificial Intelligence industry. United Nations is again urging all parties to de-escalate and return to diplomacy in the ongoing conflict between Iran and the United States in the Middle East. And following the child sexual abuse case at a Sydney childcare centre, the Australian Federal Police are appealing for information to help identify potential vict

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.sbs.com.au/language/japanese/en/podcast-episode/sbs-japanese-weekly-news-wrap-saturday-18-july/c8a9a5lzq>

137. MedFailBench: A Clinician-Built Open-Source Benchmark for Medical AI Safety Boundary Inspection

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: MedFailBench: A Clinician-Built Open-Source Benchmark for Medical AI Safety Boundary Inspection. Most medical AI benchmarks measure whether a model knows the correct answer. MedFailBench asks a different question: which safety boundary failed? We present a clinician-built synthetic benchmark and failure atlas that labels medical AI errors by severity (1--5) and safety gate type (missed urgent escalation, unsafe remote dosing, unsafe discharge reassurance, evidence fabrication, unsafe protocol ex

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.15166v1>

138. AgentKey

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AgentKey. One-stop live data marketplace for your agent. AgentKey is a plugin that connects your agent to live external data in one command. Install it into Claude Code, Codex, OpenClaw, or any MCP-based agent and instantly unlock access to search, web pages, social platforms, finance, e-commerce, business and crypto data. No integrations. No setup. Auto failover keeps workflows running.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/agentkey>

139. Beyond Data Quantity: Key Factors Driving Performance in Multilingual Language Models

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Beyond Data Quantity: Key Factors Driving Performance in Multilingual Language Models. Multilingual language models (MLLMs) are crucial for handling text across various languages, yet they often show performance disparities due to differences in resource availability and linguistic characteristics. While the impact of pre-train data percentage and model size on performance is well-known, our study reveals additional critical factors that significantly influence MLLM effectiveness.

Analyzing a wi

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2412.12500v1>

140. Comparing Physician and Artificial Intelligence Chatbot Responses to Patient Questions Posted to a Public Social Media Forum

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Comparing Physician and Artificial Intelligence Chatbot Responses to Patient Questions Posted to a Public Social Media Forum. Importance: The rapid expansion of virtual health care has caused a surge in patient messages concomitant with more work and burnout among health care professionals. Artificial intelligence (AI) assistants could potentially aid in creating answers to patient questions by drafting responses that could be reviewed by clinicians. Objective: To evaluate the ability of an AI c

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1001/jamainternmed.2023.1838>

141. Brex built its AI agent policy by watching what agents actually do, not by writing rules first

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Brex built its AI agent policy by watching what agents actually do, not by writing rules first. OpenClaw has become one of the most widely adopted agentic frameworks, but it has yet to prove itself at enterprise scale. Agents need real credentials - API keys, OAuth tokens, service accounts - to work effectively, and Brex found that traditional guardrails couldn't contain what those agents were doing with them. Brex set out to overcome these limitations by building an internal platform it calls C

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://venturebeat.com/orchestration/brex-built-its-ai-agent-policy-by-watching-what-agents-actually-do-not-by-writing-rules-first>

142. Scientific discovery in the age of artificial intelligence

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Scientific discovery in the age of artificial intelligence.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1038/s41586-023-06221-2>

143. Will A.I. Help Humanity? A Journalist Looks for Answers.

Source: New York Times (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Will A.I. Help Humanity? A Journalist Looks for Answers. In "AI for Good," Josh Tyrangiel travels to the classrooms, hospitals and research labs where people are using artificial intelligence that might benefit society. In "AI for Good," Josh Tyrangiel travels to the classrooms, hospitals and research labs where people are using artificial intelligence that might benefit society.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nytimes.com/2026/05/12/books/review/ai-for-good-josh-tyrangiel.html>

144. Measuring user competence in using artificial intelligence: validity and reliability of artificial intelligence literacy scale

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Measuring user competence in using artificial intelligence: validity and reliability of artificial intelligence literacy scale. As artificial intelligence (AI) became a part of daily life, it has become important to determine user competence in using AI technology. Here, we propose the concept of AI literacy and develop a quantitative scale for obtaining accurate data regarding the AI literacy of ordinary users. We first identified the primary core constructs of AI literacy, including awareness,

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1080/0144929x.2022.2072768>

145. COLUMN: A kid-friendly guide to AI (or for those of us born before 1975!) - Marblehead Current

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: COLUMN: A kid-friendly guide to AI (or for those of us born before 1975!) - Marblehead Current. COLUMN: A kid-friendly guide to AI (or for those of us born before 1975!) Marblehead Current

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://marbleheadcurrent.org/2026/07/15/column-a-kid-friendly-guide-to-ai-or-for-those-of-us-born-before-1975>

146. Self-Cognition in Large Language Models: An Exploratory Study

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Self-Cognition in Large Language Models: An Exploratory Study. While Large Language Models (LLMs) have achieved remarkable success across various applications, they also raise concerns regarding self-cognition. In this paper, we perform a pioneering study to explore self-cognition in LLMs. Specifically, we first construct a pool of self-cognition instruction prompts to evaluate where an LLM exhibits self-cognition and four well-designed principles to quantify LLMs' self-cognition. Our study reveals

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2407.01505v1>

147. Context Rot: Why Claude Code Sessions Decay, and How to Govern Them

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Context Rot: Why Claude Code Sessions Decay, and How to Govern Them. Long sessions rot quietly, well before any token limit is reached. Here's why, and how to govern your context in Claude Code. The post Context Rot: Why Claude Code Sessions Decay, and How to Govern Them appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/governed-context-managing-context-rot-in-claude-code>

148. Why RAG gives wrong answers (and how to fix retrieval failures)

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Why RAG gives wrong answers (and how to fix retrieval failures). In the previous post, we built a RAG system from scratch. Sixty lines of Python. Six onboarding...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/aws/why-rag-gives-wrong-answers-and-how-to-fix-retrieval-failures-gbj>

149. Osaurus

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Osaurus. Open source agents that run 100% locally on your Mac. The native macOS harness for AI agents. Any model, persistent memory, autonomous execution, cryptographic identity. Built in Swift. Fully offline. Open source.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/osaurus>

150. AI's black box problem: discovering physics we don't understand - Physics World

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI's black box problem: discovering physics we don't understand - Physics World. AI's black box problem: discovering physics we don't understand Physics World

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://physicsworld.com/a/ais-black-box-problem-discovering-physics-we-dont-understand>

151. Teaching a spreadsheet to write its own formulas

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Teaching a spreadsheet to write its own formulas. I built an AI feature into Cells, a spreadsheet app I have been building with Next.js 16, React 19,...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/aj1732/teaching-a-spreadsheet-to-write-its-own-formulas-4ah9>

152. 5 Real-World SQL Projects to Build Your Data Portfolio

Source: KDnuggets (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: 5 Real-World SQL Projects to Build Your Data Portfolio. Build a stronger data portfolio with these practical SQL projects covering customer churn, data warehousing, sales analysis, banking segmentation, and healthcare analytics.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.kdnuggets.com/5-real-world-sql-projects-to-build-your-data-portfolio>

153. Artificial Intelligence and Machine Learning in Clinical Medicine, 2023

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Artificial Intelligence and Machine Learning in Clinical Medicine, 2023. This first article in a series describes the history of artificial intelligence in medicine; the use of AI in image analysis, identification of disease outbreaks, and diagnosis; and the use of chatbots.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1056/nejmra2302038>

154. Pazi

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Pazi. Vibe code business operations. Pazi is an AI team for that idea you keep coming back to - a book, a shop, an app, a skill you want to sell. Tell Pazi what you're trying to do and it builds a team of agents around your idea and starts making things happen: a website, first outreach, content, the next step. Every time you come back, something has moved. You stay in control; your team does the rest alongside you, one real step at a time. Like vibe coding but for business operations.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/pazi>

155. Multi-Model Synthetic Training for Mission-Critical Small Language Models

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Multi-Model Synthetic Training for Mission-Critical Small Language Models. Large Language Models (LLMs) have demonstrated remarkable capabilities across many domains, yet their application to specialized fields remains constrained by the scarcity and complexity of domain-specific training data. We present a novel approach that achieves a 261x cost reduction for maritime intelligence by using LLMs as one-time teachers rather than using them directly for inference. Our method transforms 3.2 billion

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2509.13047v2>

156. Munich 1991: The Roots of the Current AI Boom

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Munich 1991: The Roots of the Current AI Boom 1991: T and P of ChatGPT, distillation, deep residual learning, LSTM, GAN When we look at the massive scale of today's Artificial Intelligence boom, it is easy to forget that the foundations of this trillion-dollar industry were laid down over 30 years ago in Munich. Today... [+34023 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://people.idsia.ch/~juergen/ai-boom-roots-munich-1991.html>

157. ARTIFICIAL INTELLIGENCE IN EDUCATION: CHALLENGES AND OPPORTUNITIES FOR SUSTAINABLE DEVELOPMENT

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: ARTIFICIAL INTELLIGENCE IN EDUCATION: CHALLENGES AND OPPORTUNITIES FOR SUSTAINABLE DEVELOPMENT. The article analyzes the role of artificial intelligence (AI) in the transformation of the educational sphere in the context of achieving the Sustainable Development Goals (SDGs). The purpose of the study is a comprehensive assessment of the potential benefits and risks of introducing AI into education in terms of their impact on sustainable development. Tasks include: identification of key AI applica

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/697ba06bfcabbbde6292d979b87b2642115f1099>

158. China's Moonshot AI Unveils Kimi Model, Threatening America's Lead

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: China's Moonshot AI Unveils Kimi Model, Threatening America's Lead China's Moonshot AI unveiled a freely available artificial intelligence model that seemed to narrow the gap with cutting-edge offerings from U.S. tech companies. The contest between China and the United States for supremacy in artificial intelligence escalated on Friday when the Chinese start-up Moonshot AI released a new A.I. model that appeared to narrow the lead held by well-funded American competitors. Mo... [1518 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nytimes.com/2026/07/17/business/china-ai-moonshot-kimi.html>

159. Fypro

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Fypro. Convert your TikTok followers into paying customers. You built a TikTok following that trusts you. Fypro turns those followers into paying customers you actually own. Drop your handle and it reads your account, then builds the whole thing for you: a site from your content, a store stocked with products matched to your niche, videos in your own voice, and a customer and email list that's yours to keep. Trained on 4M+ viral TikToks. Used by 2,000+ creators.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/fypro/launches/fypro>

160. Every AI-built site looks the same, so I built a skill that locks taste before any code is written

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Every AI-built site looks the same, so I built a skill that locks taste before any code is written. I build a lot of side projects with AI coding tools. Mostly Claude Code. A few months ago I noticed...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/codeswithroh/every-ai-built-site-looks-the-same-so-i-built-a-skill-that-locks-taste-before-any-code-is-written-4f6d>

161. New Era of Artificial Intelligence in Education: Towards a Sustainable Multifaceted Revolution

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: New Era of Artificial Intelligence in Education: Towards a Sustainable Multifaceted Revolution. The recent high performance of ChatGPT on several standardized academic tests has thrust the topic of artificial intelligence (AI) into the mainstream conversation about the future of education. As deep learning is poised to shift the teaching paradigm, it is essential to have a clear understanding of its effects on the current education system to ensure sustainable development and deployment of AI-dr

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.3390/su151612451>

162. Import AI 463: Self-improving robots; a 10k Chinese GPU cluster; and an elegiac essay for the human era

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Import AI 463: Self-improving robots; a 10k Chinese GPU cluster; and an elegiac essay for the human era. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now NVIDIA sets up a crude self-improvement loop for real world robotics:...What if you could take the best ideas from AI agents and put them into the [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/06/29/import-ai-463-self-improving-robots-a-10k-chinese-gpu-cluster-and-an-elegiac-essay-for-the-human-era>

163. From Artificial Intelligence to Explainable Artificial Intelligence in Industry 4.0: A Survey on What, How, and Where

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 60

Summary: From Artificial Intelligence to Explainable Artificial Intelligence in Industry 4.0: A Survey on What, How, and Where. Nowadays, Industry 4.0 can be considered a reality, a paradigm integrating modern technologies and innovations. Artificial intelligence (AI) can be considered the leading component of the industrial transformation enabling intelligent machines to execute tasks autonomously such as self-monitoring, interpretation, diagnosis, and analysis. AI-based methodologies (especially machine

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1109/tii.2022.3146552>

164. ARTIFICIAL INTELLIGENCE FOR THE REAL WORLD

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: ARTIFICIAL INTELLIGENCE FOR THE REAL WORLD. Artificial intelligence (A.I.) is a multidisciplinary field aimed at automating tasks that currently need human intelligence. Despite its lack of general familiarity, artificial intelligence (AI) is a technology that is revolutionizing every aspect of life. This article aims to educate laypeople about AI and encourage them to utilize it as a tool in many disciplines to rethink how we combine data, analyze it, and make choices. We quickly covered what arti

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.56726/irjmets42512>

165. eve

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: eve. Filesystem-first framework for durable backend AI agents that run anywhere.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/eve>

166. Coalition for Health AI launches implementation initiative for public health agencies - Fierce Healthcare

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Coalition for Health AI launches implementation initiative for public health agencies - Fierce Healthcare. Coalition for Health AI launches implementation initiative for public health agencies Fierce Healthcare

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.fiercehealthcare.com/ai-and-machine-learning/coalition-health-ai-launches-implementation-initiative-public-health>

167. PruneVid: Visual Token Pruning for Efficient Video Large Language Models

Source: Papers with Code (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: PruneVid: Visual Token Pruning for Efficient Video Large Language Models. In this paper, we introduce PruneVid, a visual token pruning method designed to enhance the efficiency of multi-modal video understanding. Large Language Models (LLMs) have shown promising performance in video tasks due to their extended capabilities in comprehending visual modalities. However, the substantial redundancy in video data presents significant computational challenges for LLMs. To address this issue, we introdu

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2412.16117v1>

168. Sparks of Artificial General Intelligence: Early experiments with GPT-4

Source: Semantic Scholar (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Sparks of Artificial General Intelligence: Early experiments with GPT-4. Artificial intelligence (AI) researchers have been developing and refining large language models (LLMs) that exhibit remarkable capabilities across a variety of domains and tasks, challenging our understanding of learning and cognition. The latest model developed by OpenAI, GPT-4, was trained using an unprecedented scale of compute and data. In this paper, we report on our investigation of an early version of GPT-4, when it

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.semanticscholar.org/paper/8dbd57469bb32e6d57f23f5e765bf1c9ac8e080c>

169. Don't Let Claude Grade Its Own Homework

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Don't Let Claude Grade Its Own Homework. Cross-provider PR review with Codex in GitHub Actions, and why a second opinion from a different lab beats any self-review The post Don't Let Claude Grade Its Own Homework appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/dont-let-claude-gaslight-you>

170. Ethical principles for artificial intelligence in education

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Ethical principles for artificial intelligence in education. The advancement of artificial intelligence in education (AIED) has the potential to transform the educational landscape and influence the role of all involved stakeholders. In recent years, the applications of AIED have been gradually adopted to progress our understanding of students' learning and enhance learning performance and experience. However, the adoption of AIED has led to increasing ethical risks and concerns regarding severa

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1007/s10639-022-11316-w>

171. Patreon stops asking AI bots not to scrape - and starts blocking them

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Patreon stops asking AI bots not to scrape - and starts blocking them. Patreon is strengthening its defenses against AI scraping by working with Cloudflare to block bots that train AI models on creators' content without permission. The move marks a shift away from relying on websites using robots.txt alone to actively block unauthorized AI training.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://techcrunch.com/2026/07/17/patreon-stops-asking-ai-bots-not-to-scrape-and-starts-blockin-g-them>

172. Why Your Betas Explode: The Hidden Geometry of Multicollinearity

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Why Your Betas Explode: The Hidden Geometry of Multicollinearity. Why your regression coefficients keep changing, and what geometry has to do with it. The post Why Your Betas Explode: The Hidden Geometry of Multicollinearity appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/why-your-betas-explode-the-hidden-geometry-of-multicollinearity>

173. 12 Ways to Reduce LLM Latency and Inference Costs in Production

Source: KDnuggets (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: 12 Ways to Reduce LLM Latency and Inference Costs in Production. Scaling LLMs isn't about adding GPUs. It's about removing wasted work from every request.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.kdnuggets.com/12-ways-to-reduce-llm-latency-and-inference-costs-in-production>

174. Show Me How You Reason and I'll Tell You Who You Are: Reasoning Graphs for Robust LLM Authorship Attribution

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Show Me How You Reason and I'll Tell You Who You Are: Reasoning Graphs for Robust LLM Authorship Attribution. Given the current trend to employ large language models (LLMs) in almost any imaginable context, LLM-generated text detection and authorship attribution have become a pressing issue. Prior work has primarily focused on surface-level linguistic features, an approach shown to be susceptible to paraphrasing and other obfuscation techniques. In this paper, we go beyond the linguistic surface

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14905v1>

175. Import AI 457: AI stuxnet; cursed Muon optimizer; and positive alignment

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Import AI 457: AI stuxnet; cursed Muon optimizer; and positive alignment. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now Stuxnet before Stuxnet:...Fast16 bugs software likely used in weapons programs...Here's a fascinating investigation of a ~20+ year old computer virus called fast16.sys. This software is interesting [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/05/18/import-ai-457-ai-stuxnet-cursed-muon-optimizer-and-positive-alignment>

176. Capital One releases VulnHunter, an open-source AI tool that finds software flaws before hackers do

Source: RSS Feeds (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Capital One releases VulnHunter, an open-source AI tool that finds software flaws before hackers do. Capital One on Thursday released VulnHunter, an open-source, agentic AI security tool that scans source code for exploitable vulnerabilities, maps out how an attacker would reach them, and proposes targeted fixes - all before a single line ships to production. The tool, built internally and now available on GitHub under an Apache 2.0 license, is one of the most ambitious attempts by a major financial institution.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://venturebeat.com/technology/capital-one-releases-vulnhunter-an-open-source-ai-tool-that-finds-software-flaws-before-hackers-do>

177. Ford Rehires 350 Engineers After AI Fails To Preserve Expertise or Train Juniors

Source: NewsAPI (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Ford Rehires 350 Engineers After AI Fails To Preserve Expertise or Train Juniors After Ford's automated quality-control systems and AI tools fell short, the automaker hired 350 veteran engineers over the past three years to mentor younger staff and reprogram the underperforming technology. "Artificial intelligence is a fantastic tool, but... Those engineers were "at the heart" of Ford's efforts to turn around quality problems, said Kumar Galhotra, chief operating officer. They now run mandatory training for new hires.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://slashdot.org/story/26/06/25/1624241/ford-rehires-350-engineers-after-ai-fails-to-preserve-expertise-or-train-juniors>

178. Can a Semantic Cache Become Your Primary Retrieval Layer?

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Can a Semantic Cache Become Your Primary Retrieval Layer?. Building a semantic cache layer in front of RAG - and why it might be the most underrated cost...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/surajrkhonde/can-a-semantic-cache-become-your-primary-retrieval-layer-2011>

179. Import AI 464: Fables writes GPU kernels; AI automation; and analog computation

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Import AI 464: Fables writes GPU kernels; AI automation; and analog computation. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv, cappuccinos, and feedback from readers. If you'd like to support this, please subscribe. Subscribe now Fable writes a decent GPU kernel, hinting at broader AI R&D automation:...The start of an RSI loop...Fable has written "the first genuine (and fastest) megakernel ever [...]"

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/07/06/import-ai-464-fables-writes-gpu-kernels-ai-automation-and-analog-computation>

180. Memory-Driven Self-Disclosure and Relational Turning Points: A Longitudinal Multimodal Study of Human-AI Interaction

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Memory-Driven Self-Disclosure and Relational Turning Points: A Longitudinal Multimodal Study of Human-AI Interaction. As conversational AI systems are designed for repeated use, a central question is how a series of interactions becomes a relationship. We present a longitudinal multimodal study of a memory-augmented conversational agent (24 participants x 10 sessions), in which participants rated five relational constructs -- familiarity, self-disclosure, perceived memory, conversational quality

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14593v1>

181. The Action Boundary: A Practical Safety Pattern for AI Agents

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: The Action Boundary: A Practical Safety Pattern for AI Agents. The most important line in an AI agent is not the prompt. It is the line where the program turns a...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/veobix/the-action-boundary-a-practical-safety-pattern-for-ai-agents-2loh>

182. Harnessing LLMs for Reliable Academic Supervision: A Comparative Study

Source: arXiv (research_paper)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Harnessing LLMs for Reliable Academic Supervision: A Comparative Study. Large language models routinely produce fluent answers to single-shot prompts, yet deploying them as reliable components of a domain decision system is substantially harder. Closing this gap is the work of harness engineering: the deliberate composition of deterministic scaffolding (symbolic filters, retrieval, schema-typed I/O, LLM-as-judge loops, HITL gates, persistent state, audit trails) around an LLM core. We present a

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://arxiv.org/abs/2607.14707v1>

183. Import AI 455: Automating AI Research

Source: Import AI (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Import AI 455: Automating AI Research. Welcome to Import AI, a newsletter about AI research. Import AI runs on arXiv and feedback from readers. If you'd like to support this, please subscribe. Subscribe now AI systems are about to start building themselves. What does that mean? I'm writing this post because when I look at all the publicly available information I [...]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://jack-clark.net/2026/05/04/import-ai-455-automating-ai-research>

184. Bluerails Discovery

Source: Product Hunt (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Bluerails Discovery . The rails AI agents use to find and pay you. Most "AI visibility" tools stop at telling you if AI mentions your brand. Bluerails goes further. We make you discoverable to AI agents and ready to get paid by them, on the rails we already run for marketplaces. What stands out: ? Discovery: a peer-reviewed AI-visibility score from 400 samples, not a one-off guess. Free, no signup. ?

Agent-ready checkout + global settlement ? Compliance built in Try your free Discovery report to

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.producthunt.com/products/bluerails-discovery/launches/bluerails-discovery-3>

185. Artificial intelligence in higher education: the state of the field

Source: OpenAlex (research_paper)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Artificial intelligence in higher education: the state of the field. Abstract This systematic review provides unique findings with an up-to-date examination of artificial intelligence (AI) in higher education (HE) from 2016 to 2022. Using PRISMA principles and protocol, 138 articles were identified for a full examination. Using a priori, and grounded coding, the data from the 138 articles were extracted, analyzed, and coded. The findings of this study show that in 2021 and 2022, publications rose

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://doi.org/10.1186/s41239-023-00392-8>

186. Topological Control of LLMs: A Route to Trustworthy AI - Communications of the ACM

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Topological Control of LLMs: A Route to Trustworthy AI - Communications of the ACM. Topological Control of LLMs: A Route to Trustworthy AI Communications of the ACM

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://cacm.acm.org/blogcacm/topological-control-of-llms-a-route-to-trustworthy-ai>

187. Control your ESP32 with Claude: an MCP server for your own hardware

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Control your ESP32 with Claude: an MCP server for your own hardware. Point Claude at your own IoT backend and let it read your sensors and flip your relays in plain language - no cloud middleman, no vendor skill. nodrix ships a native MCP server, owner-gated and off by default, that turns your ESP32 fleet into tools an AI agent can call.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://dev.to/decoded_cipher/control-your-esp32-with-claude-an-mcp-server-for-your-own-hardware-4ahl

188. AI Daily: Meta said in talks to lease computing power to Anthropic

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AI Daily: Meta said in talks to lease computing power to Anthropic Catch up on the top artificial intelligence news and commentary by Wall Street analysts on publicly traded companies in the space with this dail... Catch up on the top artificial intelligence news and commentary by Wall Street analysts on publicly traded companies in the space with this daily recap compiled by The Fly. TipRanks Welcomes a New ETF - NYSE:RANK TipRanks has entered a new arena in t... [3221 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://markets.businessinsider.com/news/stocks/ai-daily-meta-said-in-talks-to-lease-computing-power-to-anthropic-1036338254>

189. @microsoft/agents-a365-observability

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @microsoft/agents-a365-observability. OpenTelemetry tracing and monitoring SDK for AI agents built with TypeScript/Node.js

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@microsoft/agents-a365-observability>

190. bash-tool

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: bash-tool. Generic bash tool for AI agents, compatible with AI SDK

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/bash-tool>

191. AI coding's next CIO problem is bigger than productivity

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI coding's next CIO problem is bigger than productivity.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.informationweek.com/machine-learning-ai/the-ai-coding-rollout-worked-now-cios-have-a-bigger-problem>

192. AI-controlled F-16 begins autonomous flight testing for DARPA

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI-controlled F-16 begins autonomous flight testing for DARPA.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://militaryembedded.com/ai/machine-learning/ai-controlled-f-16-begins-autonomous-flight-testing-for-darpa>

193. Why AI leadership beats AI technology

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Why AI leadership beats AI technology.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.fastcompany.com/91574651/why-ai-leadership-beats-ai-technology>

194. What is Artificial Intelligence (AI)? | Google Cloud

Source: Serper (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: What is Artificial Intelligence (AI)? | Google Cloud. AI is a field of computer science focused on creating smart machines that can perform tasks that typically require human intelligence, like learning, reasoning ...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://cloud.google.com/learn/what-is-artificial-intelligence>

195. deepagents

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: deepagents. Deep Agents - a library for building controllable AI agents with LangGraph

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/deepagents>

196. Artificial Intelligence (AI): What it is and why it matters - SAS

Source: Serper (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Artificial Intelligence (AI): What it is and why it matters - SAS. Artificial intelligence (AI) makes it possible for machines to learn from experience, adjust to new inputs and perform human-like tasks.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://www.sas.com/en_us/insights/analytics/what-is-artificial-intelligence.html

197. @clawdbot/lobster

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 40

Summary: @clawdbot/lobster. Workflow runtime for AI agents - deterministic pipelines with approval gates

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@clawdbot/lobster>

198. EU AI Act: What Actually Applies on 2 August 2026

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: EU AI Act: What Actually Applies on 2 August 2026.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.technology.org/2026/07/17/eu-ai-act-what-actually-applies-on-2-august-2026>

199. @promptbook/azure-openai

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @promptbook/azure-openai. Promptbook: Create persistent AI agents that turn your company's scattered knowledge into action

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@promptbook/azure-openai>

200. Reddit's AI Initiative Gains Momentum: A Sign for More Upside?

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Reddit's AI Initiative Gains Momentum: A Sign for More Upside?.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://finance.yahoo.com/technology/ai/articles/reddits-ai-initiative-gains-momentum-161100148.html>

201. Applying Artificial Intelligence and machine learning in precision nutrition

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Applying Artificial Intelligence and machine learning in precision nutrition.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nature.com/articles/s41467-026-75004-w>

202. AI lab teaches itself to create superalloys

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI lab teaches itself to create superalloys.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.metaltechnews.com/story/2026/07/22/tech-bytes/ai-lab-teaches-itself-to-create-superalloys/2845.html>

203. Artificial intelligence - Wikipedia

Source: Serper (news)

Novelty: 6.0/10 | **Relevance:** 40

Summary: Artificial intelligence - Wikipedia. Artificial intelligence (AI) is the capability of computational systems to perform tasks typically associated with human intelligence

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://en.wikipedia.org/wiki/Artificial_intelligence

204. @microsoft/agents-a365-notifications

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @microsoft/agents-a365-notifications. Microsoft Agent 365 Notifications SDK for AI agents built with TypeScript/Node.js

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@microsoft/agents-a365-notifications>

205. Three Decades of Food and Drug Administration Authorizations of Artificial Intelligence/Machine Learning-Enabled Medical Devices: Persistent Specialty Concentration and the Care-Delivery Gap (1995-2025)

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Three Decades of Food and Drug Administration Authorizations of Artificial Intelligence/Machine Learning-Enabled Medical Devices: Persistent Specialty Concentration and the Care-Delivery Gap (1995-2025).

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.cureus.com/articles/500063-three-decades-of-food-and-drug-administration-authorizations-of-artificial-intelligencemachine-learning-enabled-medical-devices-persistent-specialty-concentration-and-the-care-delivery-gap-1995-2025>

206. @botonic/plugin-ai-agents

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @botonic/plugin-ai-agents. Botonic plugin for **AI agents** (e.g. OpenAI Agents-style flows) on the **current** framework line.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@botonic/plugin-ai-agents>

207. @microsoft/agents-a365-runtime

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @microsoft/agents-a365-runtime. Agent 365 Runtime SDK for AI agents built with TypeScript/Node.js

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@microsoft/agents-a365-runtime>

208. paqad-ai

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: paqad-ai. Spec-driven development framework - AI agents that think before they type

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/paqad-ai>

209. @promptbook/remote-client

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @promptbook/remote-client. Promptbook: Create persistent AI agents that turn your company's scattered knowledge into action

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@promptbook/remote-client>

210. AMD Ryzen 7 7700X3D Review - Machine Learning / AI

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AMD Ryzen 7 7700X3D Review - Machine Learning / AI.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.techpowerup.com/review/amd-ryzen-7-7700x3d/5.html>

211. Vantage AI Announces \$100M+ Capital Initiative to Advance AI-Driven Algorithmic Technology

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Vantage AI Announces \$100M+ Capital Initiative to Advance AI-Driven Algorithmic Technology.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.streetinsider.com/Newsfile/Vantage+AI+Announces+%24100M%2B+Capital+Initiative+to+Advance+AI-Driven+Algorithmic+Technology/26785526.html>

212. Google AI - How we're making AI helpful for everyone

Source: Serper (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Google AI - How we're making AI helpful for everyone. Discover how Google AI is committed to enriching knowledge, solving complex challenges and helping people grow by building useful AI tools and technologies.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://ai.google/>

213. @microsoft/agents-a365-tooling

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @microsoft/agents-a365-tooling. Agent 365 Tooling SDK for AI agents built with TypeScript/Node.js

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@microsoft/agents-a365-tooling>

214. @promptbook/utis

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @promptbook/utis. Promptbook: Create persistent AI agents that turn your company's scattered knowledge into action

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@promptbook/utis>

215. How artificial intelligence is changing the mental health space

Source: SerpAPI News (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: How artificial intelligence is changing the mental health space.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://www.purdueexponent.org/lifestyles/health/how-artificial-intelligence-is-changing-the-mental-health-space/article_709a6493-1cab-5e8e-b503-1ca0ec7df5f8.html

216. @anthropic-ai/claude-agent-sdk

Source: npm (repo)

Novelty: 6.0/10 | **Relevance:** 20

Summary: @anthropic-ai/claude-agent-sdk. SDK for building AI agents with Claude Code's capabilities. Programmatically interact with Claude to build autonomous agents that can understand codebases, edit files, and execute workflows.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.npmjs.com/package/@anthropic-ai/claude-agent-sdk>

217. AI-powered target identification: How machine learning is finding new drug targets - Drug Discovery News

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: AI-powered target identification: How machine learning is finding new drug targets - Drug Discovery News. AI-powered target identification: How machine learning is finding new drug targets Drug Discovery News

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.drugdiscoverynews.com/ai-powered-target-identification-how-machine-learning-is-finding-new-drug-targets-17344>

218. I built an open-source workspace to tame my AI coding agents

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: I built an open-source workspace to tame my AI coding agents. The problem Every AI coding CLI is a snowflake: different flags, different session models,...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: https://dev.to/arnab_karmakar_41f6ef02ca/i-built-an-open-source-workspace-to-tame-my-ai-coding-agents-pb2

219. AI Help With Your Database Design, Without the Cloud

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: AI Help With Your Database Design, Without the Cloud. A BYOK ERD tool changes the deal: private AI database design with your own OpenAI, Claude, Gemini, or DeepSeek key. AI ERD without cloud accounts, inside...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/tbson87/ai-help-with-your-database-design-without-the-cloud-3214>

220. LOOM Day 13 - from "what the code does" to "how much"

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: LOOM Day 13 - from "what the code does" to "how much". LOOM is a tiny effect-typed language: a trust layer for AI-written code. Every function must honestly...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/umbraaeternaa/loom-day-13-from-what-the-code-does-to-how-much-263m>

221. China's Kimi K3 rattles US AI industry

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: China's Kimi K3 rattles US AI industry WASHINGTON - Kimi K3, a new artificial intelligence program from a Chinese startup called Moonshot AI, stunned the US tech industry on Friday, setting off fresh discussion over the China-US rivalry to dominate AI. WASHINGTON - Kimi K3, a new artificial intelligence program from a Chinese startup called Moonshot AI, stunned the US tech industry on Friday, setting off fresh discussion over the China-US rivalry to dominate AI. The program was released Thursday

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.manilatimes.net/2026/07/18/world/chinas-kimi-k3-rattles-us-ai-industry/2386827>

222. The Change Must Defend Itself: A Five-R Review for AI-Generated Code

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: The Change Must Defend Itself: A Five-R Review for AI-Generated Code. How to evaluate reason, reach, reversibility, resilience and reviewability before an AI-assisted...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://dev.to/karanratra07/the-change-must-defend-itself-a-five-r-review-for-ai-generated-code-3eep>

223. How AiClaw Hardens Local Agent Runtimes on Your Machine

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: How AiClaw Hardens Local Agent Runtimes on Your Machine. AiClaw's latest local-runtime work fixes a practical problem that shows up fast in self-hosted agent...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/chowyu12/how-aiclaw-hardens-local-agent-runtimes-on-your-machine-1nkc>

224. # Your AI Agent Is a Distributed System in Disguise - Instrument It with OpenTelemetry and SigNoz Before It Bankrupts You

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: # Your AI Agent Is a Distributed System in Disguise - Instrument It with OpenTelemetry and SigNoz Before It Bankrupts You. My warm-up blog for the Agents of SigNoz hackathon by WeMakeDevs x SigNoz. Here's a claim I'll...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/sairaammmbharadwaj/-your-ai-agent-is-a-distributed-system-in-disguise-instrument-it-with-opentelemetry-and-signoz-2e1d>

225. Morgan State University Launches Artificial Intelligence Degree to Prepare Students for the Next Era of Innovation - Morgan State University

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Morgan State University Launches Artificial Intelligence Degree to Prepare Students for the Next Era of Innovation - Morgan State University. Morgan State University Launches Artificial Intelligence Degree to Prepare Students for the Next Era of Innovation Morgan State University

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.morgan.edu/news/morgan-launches-artificial-intelligence-degree>

226. Using Classical ML to Empower AI Agents - towardsdatascience.com

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Using Classical ML to Empower AI Agents - towardsdatascience.com. Using Classical ML to Empower AI Agents towardsdatascience.com

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/using-classical-ml-to-empower-ai-agents>

227. The Dawn of AI Change-Control - Communications of the ACM

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 0

Summary: The Dawn of AI Change-Control - Communications of the ACM. The Dawn of AI Change-Control Communications of the ACM

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://cacm.acm.org/news/the-dawn-of-ai-change-control>

228. Deep Reinforcement Learning-based combat recognition of traditional Chinese Sanda under artificial intelligence technology - Nature

Source: Google News RSS (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Deep Reinforcement Learning-based combat recognition of traditional Chinese Sanda under artificial intelligence technology - Nature. Deep Reinforcement Learning-based combat recognition of traditional Chinese Sanda under artificial intelligence technology Nature

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.nature.com/articles/s41598-026-62617-w>

229. Instrumenting an AI-Powered GitHub Analyzer with OpenTelemetry and SigNoz

Source: Dev.to (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Instrumenting an AI-Powered GitHub Analyzer with OpenTelemetry and SigNoz. This article is my submission for the Agents of SigNoz Hackathon: Blog Track, where participants...

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://dev.to/divyasinghdev/instrumenting-an-ai-powered-github-analyzer-with-opentelemetry-and-signoz-55l3>

230. Chinese AI model takes U.S. tech industry by surprise

Source: GNews (news)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Chinese AI model takes U.S. tech industry by surprise Another powerful new artificial intelligence model from China is taking the U.S. tech industry by surprise. Another powerful new artificial intelligence model from China took the U.S. tech industry by surprise Friday, the latest sign that Chinese startups that publicly release their "open-source" AI technology are making the California titans of AI sweat. ... [4756 chars]

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL:

<https://www.nbcchicago.com/news/tech/chinese-ai-model-kimi-k3-rivaling-claude-chatgpt/3963114>

231. How to Work Effectively with GPT-5.6

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: How to Work Effectively with GPT-5.6. Maximize the latest OpenAI model The post How to Work Effectively with GPT-5.6 appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/how-to-work-effectively-with-gpt-5-6>

232. Analog AI Is Back, But Can It Survive Its Own Noise?

Source: Towards Data Science (blog)

Novelty: 6.0/10 | **Relevance:** 0

Summary: Analog AI Is Back, But Can It Survive Its Own Noise?. AI's energy crisis is reviving an old idea: computing with physics instead of digital logic. Here's how analog chips actually work, why noise nearly killed the idea once already, and what happens when you simulate that noise yourself. The post Analog AI Is Back, But Can It Survive Its Own Noise? appeared first on Towards Data Science .

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://towardsdatascience.com/analog-ai-is-back-can-it-survive-its-own-noise>

233. Git Worktrees for AI Development

Source: KDnuggets (blog)

Novelty: 6.0/10 | **Relevance:** 20

Summary: Git Worktrees for AI Development. A Git worktree is a separate directory checked out from the same repository. You can have as many as you need, each on its own branch, all coexisting simultaneously on your filesystem.

Recommended Action: Candidate evidence for enterprise AI strategy and implementation decisions.

URL: <https://www.kdnuggets.com/git-worktrees-for-ai-development>